

# Assignment 06 - Submitted By

Muhammad Raaid Khan

DS & AI

Batch-05

## Task

Building **ResNet152V2** Model for Image Classification with Small Dataset

**Number of classes:** 20 (Classes 0-19) **Classes** = owl | galaxy | lightning | wine-bottle | t-shirt | waterfall | sword | school-bus | calculator | sheet-music | airplanes | lightbulb | skyscraper | mountain-bike | fireworks | computer-monitor | bear | grand-piano | kangaroo | laptop ]

**Dataset Structure** Two folders: Training: 1554 images Test: 500 images

**Dataset Link** <https://github.com/miladfa7/Image-Classification-Transfer-Learning/tree/master/Dataset>

## Suppressing Tensorflow Warnings

```
In [150... import os
import warnings
warnings.filterwarnings('ignore')

# Suppress TensorFlow warnings
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
```

## Using On-Machine GPU

```
In [173... from tensorflow.python.client import device_lib

# List physical GPU devices
gpus = tf.config.list_physical_devices('GPU')

if gpus:
    print(f"Number of GPUs available: {len(gpus)}\n")
```

```

for i, gpu in enumerate(gpus):
    print(f"GPU {i}:")
    print(f"Name: {gpu.name}")
    print(f"Device Type: {gpu.device_type}")

    # Get detailed device info
    gpu_details = device_lib.list_local_devices()
    for device in gpu_details:
        if device.device_type == 'GPU':
            print(f"Device Description: {device.physical_device_desc}")
            print(f"Memory Limit: {device.memory_limit / (1024 ** 3):.2f} GB\n")
else:
    print("No GPU available")

```

Number of GPUs available: 1

GPU 0:

Name: /physical\_device:GPU:0

Device Type: GPU

Device Description: device: 0, name: NVIDIA GeForce RTX 3060 Laptop GPU, pci bus id: 0000:01:00.0, compute capability: 8.6

Memory Limit: 3.50 GB

## Importing Necessary Libraries

```

In [151... import tensorflow as tf
import numpy as np
import os

```

```

In [152... import zipfile
zip = zipfile.ZipFile("/tf/Image-Classification-Transfer-Learning/Dataset.zip", 'r')
zip.extractall("/tf/Image-Classification-Transfer-Learning/Dataset")
zip.close()

```

```

In [153... dataset_dir = os.path.join('/tf/Image-Classification-Transfer-Learning/Dataset/')

```

```

In [154... from tensorflow.keras.preprocessing.image import ImageDataGenerator

```

```

data_gen = ImageDataGenerator(rotation_range=50,
                               width_shift_range=0.2,
                               height_shift_range=0.2,
                               zoom_range=0.3,
                               horizontal_flip=True,
                               vertical_flip=True,

```



```

class_mode='categorical',
shuffle=True,
subset='training', # Use only the training portion
seed=SEED) # targets are directly converted into one-hot vectors

# Validation
valid_gen = data_gen.flow_from_directory(training_dir,
                                         target_size=(256, 256),
                                         batch_size=Batch_size,
                                         classes=classes,
                                         class_mode='categorical',
                                         shuffle=False,
                                         subset='validation', # Use the validation portion
                                         seed=SEED)

```

Found 1247 images belonging to 20 classes.  
Found 307 images belonging to 20 classes.

## Visualizing Data

```

In [158... CLASS_NAMES = np.array(['owl', 'galaxy', 'lightning', 'wine-bottle', 't-shirt', 'waterfall', 'sword', 'school-bus',
                             'calculator', 'sheet-music', 'airplanes', 'lightbulb', 'skyscraper', 'mountain-bike', 'fireworks',
                             'computer-monitor', 'bear', 'grand-piano', 'kangaroo', 'laptop', ], dtype='<U10')

import matplotlib.pyplot as plt

def show_batch(image_batch, label_batch):
    plt.figure(figsize=(25,20))
    for n in range(8):
        ax = plt.subplot(1,8,n+1)
        plt.imshow(image_batch[n])
        plt.title(CLASS_NAMES[label_batch[n]==1][0].title())
        plt.axis('off')

image_batch, label_batch = next(train_gen)
show_batch(image_batch, label_batch)

```



# Getting ResNet Model

```
In [159... ResNet_model = tf.keras.applications.ResNet152V2(weights='imagenet', include_top=False, input_shape=(img_h, img_w, 3))
```

## Building Model

```
In [175... from tensorflow.keras import Model
from tensorflow.keras.layers import Conv2D, Dense, MaxPooling2D, Dropout, Flatten, GlobalAveragePooling2D
from tensorflow.keras.models import Sequential

# The last 15 layers fine tune
for layer in ResNet_model.layers[:15]:
    layer.trainable = False

x = ResNet_model.output
x = GlobalAveragePooling2D()(x)
x = Flatten()(x)
x = Dense(units=512, activation='relu')(x)
x = Dropout(0.3)(x)
x = Dense(units=512, activation='relu')(x)
x = Dropout(0.3)(x)
output = Dense(units=20, activation='softmax')(x)
model = Model(ResNet_model.input, output)
```

```
In [161... loss = tf.keras.losses.CategoricalCrossentropy()
optimizer = tf.keras.optimizers.Adam(learning_rate=0.001)
model.compile(optimizer=optimizer, loss=loss, metrics= ['accuracy'])
```

```
In [162... STEP_SIZE_TRAIN=train_gen.n//train_gen.batch_size
STEP_SIZE_VALID=valid_gen.n//valid_gen.batch_size
transfer_learning_history = model.fit_generator(generator=train_gen,
        steps_per_epoch=STEP_SIZE_TRAIN,
        validation_data=valid_gen,
        validation_steps=STEP_SIZE_VALID,
        epochs=20,
    )
```

Epoch 1/20  
155/155 [=====] - 28s 116ms/step - loss: 1.5425 - accuracy: 0.5811 - val\_loss: 3.7060 - val\_accuracy: 0.7237  
Epoch 2/20  
155/155 [=====] - 15s 98ms/step - loss: 1.0507 - accuracy: 0.7417 - val\_loss: 1.0842 - val\_accuracy: 0.7763  
Epoch 3/20  
155/155 [=====] - 15s 99ms/step - loss: 0.7430 - accuracy: 0.8111 - val\_loss: 0.8066 - val\_accuracy: 0.8059  
Epoch 4/20  
155/155 [=====] - 15s 98ms/step - loss: 0.6485 - accuracy: 0.8313 - val\_loss: 0.6108 - val\_accuracy: 0.8355  
Epoch 5/20  
155/155 [=====] - 16s 100ms/step - loss: 0.5288 - accuracy: 0.8499 - val\_loss: 0.6217 - val\_accuracy: 0.8520  
Epoch 6/20  
155/155 [=====] - 17s 108ms/step - loss: 0.4496 - accuracy: 0.8757 - val\_loss: 0.6264 - val\_accuracy: 0.8553  
Epoch 7/20  
155/155 [=====] - 17s 110ms/step - loss: 0.4161 - accuracy: 0.8862 - val\_loss: 0.3740 - val\_accuracy: 0.8914  
Epoch 8/20  
155/155 [=====] - 16s 103ms/step - loss: 0.3272 - accuracy: 0.9080 - val\_loss: 0.5937 - val\_accuracy: 0.8783  
Epoch 9/20  
155/155 [=====] - 17s 107ms/step - loss: 0.3226 - accuracy: 0.9080 - val\_loss: 0.5047 - val\_accuracy: 0.8684  
Epoch 10/20  
155/155 [=====] - 16s 104ms/step - loss: 0.3878 - accuracy: 0.8991 - val\_loss: 1.1132 - val\_accuracy: 0.8388  
Epoch 11/20  
155/155 [=====] - 17s 107ms/step - loss: 0.3211 - accuracy: 0.9048 - val\_loss: 0.3937 - val\_accuracy: 0.9013  
Epoch 12/20  
155/155 [=====] - 17s 110ms/step - loss: 0.3151 - accuracy: 0.9112 - val\_loss: 0.5092 - val\_accuracy: 0.8553  
Epoch 13/20  
155/155 [=====] - 17s 110ms/step - loss: 0.3556 - accuracy: 0.9040 - val\_loss: 0.5705 - val\_accuracy: 0.8520  
Epoch 14/20  
155/155 [=====] - 17s 106ms/step - loss: 0.2845 - accuracy: 0.9266 - val\_loss: 0.4730 - val\_accuracy: 0.8882  
Epoch 15/20  
155/155 [=====] - 16s 102ms/step - loss: 0.2562 - accuracy: 0.9241 - val\_loss: 0.7709 - val\_accuracy: 0.8586  
Epoch 16/20

```
155/155 [=====] - 16s 105ms/step - loss: 0.2622 - accuracy: 0.9282 - val_loss: 0.4274 - val_accuracy: 0.8980
Epoch 17/20
155/155 [=====] - 16s 100ms/step - loss: 0.2528 - accuracy: 0.9298 - val_loss: 0.4509 - val_accuracy: 0.8816
Epoch 18/20
155/155 [=====] - 15s 99ms/step - loss: 0.3066 - accuracy: 0.9144 - val_loss: 0.4050 - val_accuracy: 0.8947
Epoch 19/20
155/155 [=====] - 16s 100ms/step - loss: 0.2453 - accuracy: 0.9338 - val_loss: 0.7002 - val_accuracy: 0.8553
Epoch 20/20
155/155 [=====] - 16s 105ms/step - loss: 0.1634 - accuracy: 0.9500 - val_loss: 0.3648 - val_accuracy: 0.9145
```

## Visualizing Model Accuracy and Loss

```
In [163... import matplotlib.pyplot as plt


acc = transfer_learning_history.history['accuracy']
val_acc = transfer_learning_history.history['val_accuracy']

loss = transfer_learning_history.history['loss']
val_loss = transfer_learning_history.history['val_loss']

epochs_range = range(20)

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(epochs_range, acc, label='Training Accuracy')
plt.plot(epochs_range, val_acc, label='Validation Accuracy')
plt.legend(loc='lower right')
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(epochs_range, loss, label='Training Loss')
plt.plot(epochs_range, val_loss, label='Validation Loss')
plt.legend(loc='upper right')
plt.title('Training and Validation Loss')
plt.show()
```



```
In [164... model.evaluate(valid_gen, steps=STEP_SIZE_VALID, verbose=1)
```

```
38/38 [=====] - 3s 83ms/step - loss: 0.4971 - accuracy: 0.8750
```

```
Out[164... [0.49708300828933716, 0.875]
```



## Making Predictions on Un-Seen Images

```
In [165... import cv2
from tensorflow.keras.preprocessing import image

# Path to the new images
new_image_path = '/tf/Test Pics/fw.jpg'

# Load the image
img = image.load_img(new_image_path, target_size=(256, 256)) # Ensure the image is resized to the model's expected input size

# Preprocess the image (convert to array and scale it)
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension (1, height, width, channels)
img_array = img_array / 255.0 # Normalize the pixel values (same as you did during training)

test_img = cv2.imread(new_image_path)
plt.imshow(test_img)
```

```
Out[165... <matplotlib.image.AxesImage at 0x7f53489fe950>
```



```
In [166... # Make predictions
predictions = model.predict(img_array)

# Get the predicted class index
predicted_class_index = np.argmax(predictions, axis=1)

# If you have class names, map the index to the corresponding class label
class_names = classes # Replace with your actual class names
predicted_class = class_names[predicted_class_index[0]]

# Output the predicted class
print(f"Predicted class: {predicted_class}")

1/1 [=====] - 2s 2s/step
1/1 [=====] - 2s 2s/step
Predicted class: fireworks
```

```
In [169... import cv2
from tensorflow.keras.preprocessing import image


# Path to the new images
new_image_path = '/tf/Test Pics/laptop.png'

# Load the image
img = image.load_img(new_image_path, target_size=(256, 256)) # Ensure the image is resized to the model's expected input size

# Preprocess the image (convert to array and scale it)
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension (1, height, width, channels)
img_array = img_array / 255.0 # Normalize the pixel values (same as you did during training)

test_img = cv2.imread(new_image_path)
plt.imshow(test_img)
```

```
Out[169... <matplotlib.image.AxesImage at 0x7f5348a0e810>
```



```
In [170... # Make predictions
predictions = model.predict(img_array)

# Get the predicted class index
predicted_class_index = np.argmax(predictions, axis=1)

# If you have class names, map the index to the corresponding class label
class_names = classes # Replace with your actual class names
predicted_class = class_names[predicted_class_index[0]]

# Output the predicted class
print(f"Predicted class: {predicted_class}")

1/1 [=====] - 0s 62ms/step
1/1 [=====] - 0s 62ms/step
Predicted class: laptop
```

```
In [171... import cv2
from tensorflow.keras.preprocessing import image


# Path to the new images
new_image_path = '/tf/Test Pics/bulb.png'
```

```
# Load the image
img = image.load_img(new_image_path, target_size=(256, 256)) # Ensure the image is resized to the model's expected input size

# Preprocess the image (convert to array and scale it)
img_array = image.img_to_array(img)
img_array = np.expand_dims(img_array, axis=0) # Add batch dimension (1, height, width, channels)
img_array = img_array / 255.0 # Normalize the pixel values (same as you did during training)

test_img = cv2.imread(new_image_path)
plt.imshow(test_img)
```

Out[171...] <matplotlib.image.AxesImage at 0x7f5348be9250>



```
In [172...] # Make predictions
predictions = model.predict(img_array)

# Get the predicted class index
predicted_class_index = np.argmax(predictions, axis=1)

# If you have class names, map the index to the corresponding class label
```

```
class_names = classes # Replace with your actual class names
predicted_class = class_names[predicted_class_index[0]]

# Output the predicted class
print(f"Predicted class: {predicted_class}")
```

1/1 [=====] - 0s 49ms/step

1/1 [=====] - 0s 49ms/step

Predicted class: lightbulb